


COMP4431 Artificial Intelligence

Knowledge based Agent and Expert System

Raymond Pang
Department of Computing
The Hong Kong Polytechnic University





Agenda

- Knowledge based agent (KBA)
 - Knowledge base and representation
- Inference Engine
 - Forward Chaining
 - Backward Chaining
- Game
 - MiniMax Algorithm
 - Alpha-beta Pruning

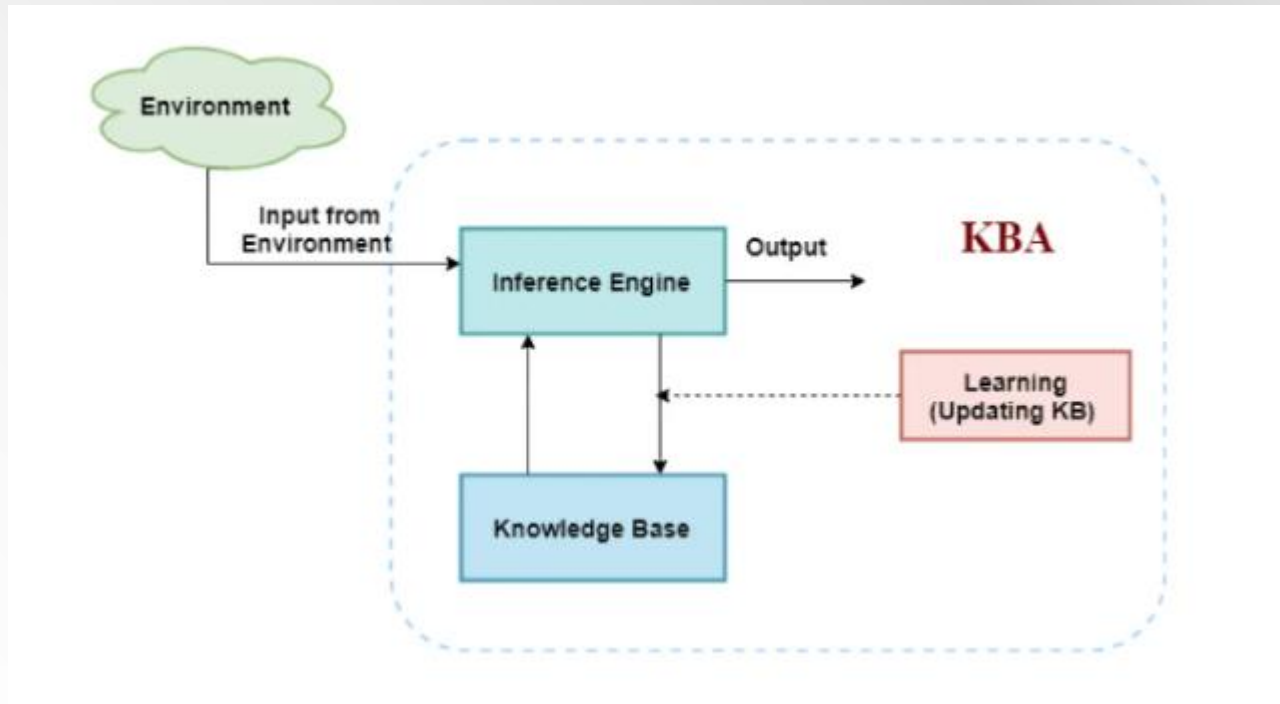


Knowledge-based Agent

- An Intelligent agent needs knowledge about the world for taking decisions and reasoning to act efficiently
- Knowledge-based agents (KBA) are those agents who have the capability of
 - ❑ maintaining an internal state of knowledge,
 - ❑ reason over that knowledge,
 - ❑ update their knowledge after observations and,
 - ❑ take action.
- These agents can represent the world with some formal representation and act intelligently

Knowledge-based Agent

- KBAs are composed of 2 main parts
 - Knowledge-base
 - Inference Engine





Knowledge-based Agent

- Knowledge base: a central component/repository of all knowledge in the KBA
 - A collection of sentences (here ‘sentence’ is a technical term and it is not identical to sentence in natural language like English)
 - These sentences are expressed in so-called a knowledge representation language
- In practice, the knowledge bases are built based on human/expert experience
 - It can also be learned by the agent itself, but it will be based on other more complex techniques like machine learning or reinforcement learning, etc.

Inference Rules

- One common way of knowledge base representation
 - Sometimes referring as Production Rules
- A production rule consists of (condition, decision) pairs which mean “if condition then decision”
 - Here “decision” can also be the condition for another rule, not necessary to be a real action to perform
- We said a rule is fired if its condition is fulfilled




Inference Rules

- A KB can have many rules, e.g.
 - IF(at bus stop AND bus arrives) THEN decision (get into the bus)
 - IF(on the bus AND paid AND empty seat) THEN decision (sit down)
 - IF (on bus AND unpaid) THEN decision (pay charges)
 - IF (on the bus AND sit down) THEN decision (play mobile phone)
 - IF (bus arrives at destination) THEN decision (get down from the bus)
- As one rule is fired, it may trigger another rule to be fired. There can be a chain reaction (referred as **inference chain**)!

Forward Chaining

- Forward chaining is a method of reasoning in the Inference Engine
- Inference rules are applied to existing data to extract additional data until an endpoint (goal) is achieved.
- This approach is often used in expert systems for tasks such as troubleshooting and diagnostics.





Forward Chaining Steps

- The system is given one or more facts (from the environment or user input)
- Then the rules are searched in the KB for each fact. Rules that fulfil the conditions are selected (i.e. IF part)
- Now each rule is able to produce new conditions from the conclusion of the invoked one.
- The added conditions are processed again by repeating step 2. The process will end if there is no new conditions exist.



Example of Forward Chaining: A simple weather forecast system

- The system starts with KB :

- Rule 1

- IF the ambient temperature is above 90°F

- THEN the weather is hot

- Rule 2

- IF the relative humidity is greater than 65%

- THEN the atmosphere is humid

- Rule 3

- IF the weather is hot and the atmosphere is humid

- THEN thunderstorms are likely to develop



Example of Forward Chaining: A simple weather forecast system

From:

- Fact 1 (Input)
the ambient temperature is 92°F
- Rule 1
IF the ambient temperature is above 90°F
THEN the weather is hot
- The inference engine can deduce:
the weather is hot



Example of Forward Chaining: A simple weather forecast system

From:

- Fact 2 (Input)
the weather is hot
- Fact 3 (Input)
the atmosphere is humid
- Rule 3
IF the weather is hot and the atmosphere is humid
THEN thunderstorms are likely to develop
- The inference engine can deduce:
thunderstorms are likely to develop



Properties of Forward Chaining

- **Data-Driven:** The reasoning starts from available data (facts) and works toward a goal.
- **Bottom-Up Approach:** It builds knowledge from facts, gradually moving towards conclusions.
- **Breadth-First Search Strategy:** The inference engine explores multiple rules simultaneously, applying them step by step.
- **Possibility of Irrelevant Rules:** Forward chaining may explore rules that do not contribute to the final solution, making it less efficient in some cases.

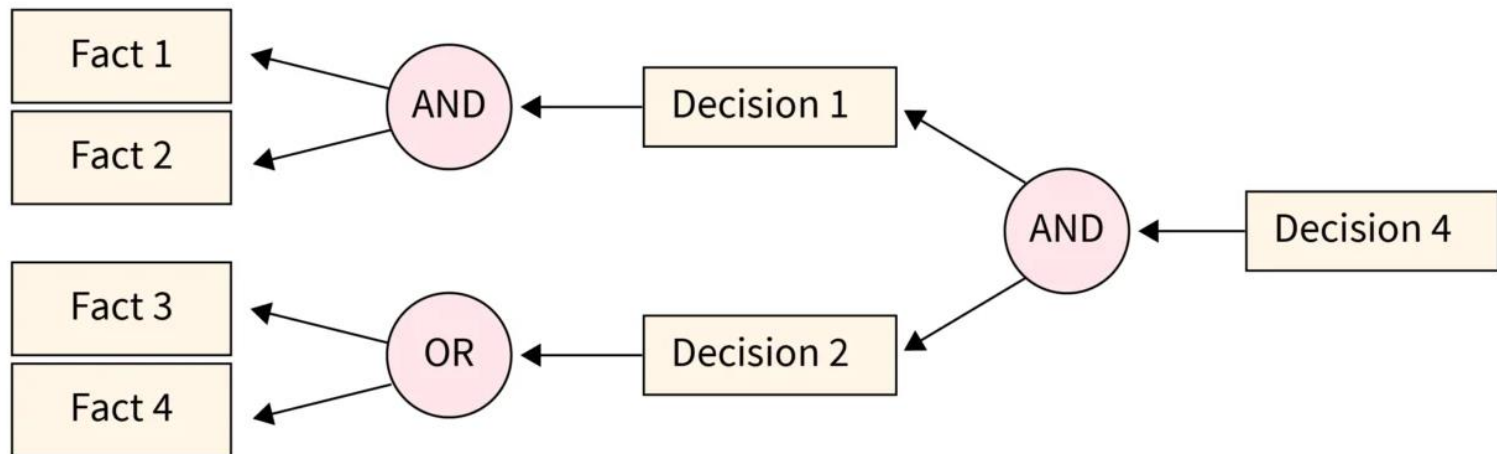
Ways of Reasoning

- Forward chaining / reasoning
 - Start with all the known data and progress toward the conclusion
- An alternative: we select a possible conclusion and try to prove its validity by looking for supporting evidence
 - It's like what a detective does!



Backward Chaining

- Backward chaining is a goal-driven reasoning strategy for inference.
- It starts with a goal and works backward to determine if the available facts support the goal.



Example of Backward Chaining

- Goal:

- ☐ Does the patient have the flu?

- Rule:

- ☐ IF the patient has a fever and sore throat, THEN they might have the flu.

- Sub-goals:

- ☐ Verify if the patient has a fever.

- ☐ Verify if the patient has a sore throat.



Backward Chaining

- **Goal-Driven:** Reasoning begins with a desired goal and searches for evidence to support it.
- **Top-Down Approach:** The system starts from the goal and works back to find relevant facts.
- **Depth-First Search Strategy:** The inference engine follows a path deeply before exploring other possibilities, prioritizing each goal or sub-goal in sequence.
- **Possibility of Infinite Loops:** If not handled properly, backward chaining may get stuck in loops while looking for evidence to support the goal.



A Simple Medical Diagnosis System

- Rule 1

IF (nasal congestion and
viremia)

THEN influenza

- Rule 2

IF (runny nose)

THEN nasal congestion

- Rule 3

IF (body-aches > 100)

THEN achiness

- Rule 4

IF (temperature > 100)

THEN fever

- Rule 5

IF (headache)

THEN achiness

- Rule 6

IF (fever and
achiness and
cough)

THEN viremia

Forward Chaining Results

- For Input :

- ☐ Runny nose

- ☐ Temperature = 101.7

- ☐ Headache

- ☐ Cough


- Trace of rule execution

Execution Cycle	Premises	Selected Rule	Derived Fact
1	runny nose	Rule 2	nassal congestion
2	temperature	Rule 4	fever
3	headache	Rule 5	achiness
4	fever, achiness, cough	Rule 6	viremia
5	nassal congestion, viremia	Rule 1	influenza

Backward Chaining Results

- Initial goal : diagnosis influenza
- Trace of rule execution

Execution Cycle	Goal resolving	Rule fired	Goals created & accumulated	Premise satisfied
1	influenza	Rule 1	nassal congestion, viremia	nil
2	nassal congestion	Rule 2	runny nose, viremia	runny nose
3	viremia	Rule 6	fever, achiness, cough	cough
4	fever	Rule 4	Temperature > 100, achiness	Temperature > 100
5	achiness	Rule 3	Body-aches, achiness	nil
5	achiness	Rule 5	headache	headache
6	All goals are cleared and satisfied premises, so this hypothesis was correct, subject has influenza			



Game



Games

- Multiagent environment

- ☐ In which each agent needs to consider the actions of other agents
- ☐ And how they affect its own welfare

- Games

- ☐ Views multiagent environment as a game, provided that the impact of each agent on the other is significant

- Games in AI

- ☐ The state of a game is easy to represent
- ☐ Agents are usually restricted to a small number of actions
- ☐ Outcomes are defined by precise rules



Game Playing Strategy

- Maximize winning possibility assuming that opponent will try to minimize (Minimax Algorithm)
- Ignore the unwanted portion of the search tree (Alpha Beta Pruning)
- Evaluation(Utility) Function
 - A measure of winning possibility of the player

Formulate Game to Search Problem

- S_0
 - The initial state, which specifies how the game is set up at the start
- $\text{Player}(s)$
 - Defines which players has the move in a state
- $\text{Actions}(s)$
 - Defines a set of legal move in a state
- $\text{Result}(s, a)$
 - The transition model, which defines the result of a move
- $\text{Terminal-Test}(s)$
 - A terminal test, which is true when the game is over
- $\text{Utility}(s, p)$
 - A utility function, also called an objective function or payoff function
 - Defines the final numeric value for a game that ends in terminal state s for a player p

Tic-Tac-Toe

X	O	

$$e(p) = 6 - 5 = 1$$

- ▶ **Initial State:** Board position of 3x3 matrix with 0 and X.
- ▶ **Actions:** Putting 0's or X's in vacant positions alternatively
- ▶ **Terminal test:** Which determines game is over
- ▶ **Utility function:**


$$e(p) = (\text{No. of complete rows, columns or diagonals are still open for player}) - (\text{No. of complete rows, columns or diagonals are still open for opponent})$$



Minimax Algorithm

- Generate the game tree
- Apply the utility function to each terminal state to get its value
- Use these values to determine the utility of the nodes one level higher up in the search tree
 - From bottom to top
 - For a max level, select the maximum value of its successors
 - For a min level, select the minimum value of its successors
- From root node select the move which leads to highest value

Game tree for Tic-Tac-Toe



MAX

Maximizing O
Best move: [1, 1] (center)

Depth = 3

O	O	X
X		
	O	X

MIN

O	O	X
X	O	
	O	X

+1

O	O	X
X		O
	O	X

0

O	O	X
X		
O	O	X

-1

Depth = 2

MAX

Depth = 1

+1

O	O	X
X		O
X	O	X

0

O	O	X
X	X	O
	O	X

O	O	X
X	X	
O	O	X

0

O	O	X
X		X
O	O	X

-1

MIN

Depth = 0

+1

O	O	X
X	O	O
X	O	X

0

O	O	X
X	X	O
O	O	X

O	O	X
X	X	O
O	O	X

0



Properties of Minimax

- **Complete** : Yes (if tree is finite)
- **Time complexity** : $O(b^d)$
- **Space complexity** : $O(bd)$ (depth-first exploration)



Observation

- Minimax algorithm, presented above, requires expanding the entire state-space.
- Severe limitation, especially for problems with a large state-space.
- Some nodes in the search can be *proven to be irrelevant to the outcome* of the search

Alpha-Beta Strategy

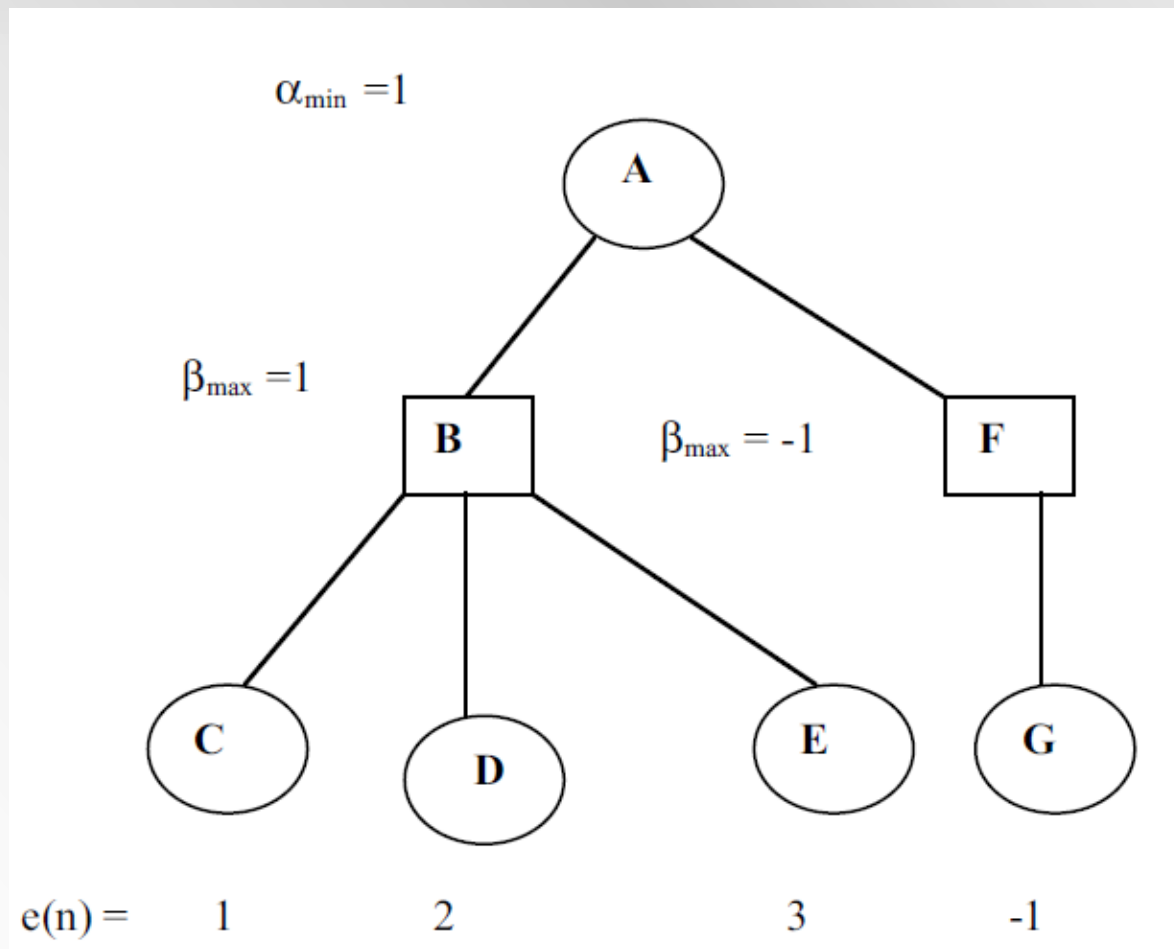
- Maintain two bounds:
 - Alpha (α): a lower bound on best that the player to move can achieve
 - Beta (β): an upper bound on what the opponent can achieve
- Search, maintaining α and β
- Whenever $\alpha \geq \beta_{\text{higher}}$, or $\beta \leq \alpha_{\text{higher}}$ further search at this node is irrelevant



How to Prune the Unnecessary Path

- **If** beta value of any MIN node below a MAX node is less than or equal to its alpha value, **then** prune the path below the MIN node.
- **If** alpha value of any MAX node below a MIN node exceeds the beta value of the MIN node, **then** prune the nodes below the MAX node.


Example



Tic-Tac-Toe

(MAX) Start


$$e(p) = 0$$



X : MAX player

O : MIN player

$e(p) = (\text{rows} + \text{cols} + \text{diagonals open to 'X'}) - (\text{Same to 'O'})$




X's Turn

$$e = 8 - 4 = 4$$

$$e = 8 - 5 = 3$$

$$e = 8 - 6 = 2$$



O's Turn

$$e = 5 - 4 = 1$$

$$e = 5 - 3 = 2$$

Alpha-Beta Search Algorithm

- If the MAXIMIZER nodes already possess α_{\min} values, then their current **α_{\min} value = Max (α_{\min} value, α'_{\min})**; on the other hand, if the MINIMIZER nodes already possess β_{\max} values, then their current **β_{\max} value = Min (β_{\max} value, β'_{\max})**.
- If the estimated β_{\max} value of a MINIMIZER node N is less than the α_{\min} value of its parent MAXIMIZER node N' then there is no need to search below the node MINIMIZER node N. Similarly, if the α_{\min} value of a MAXIMIZER node N is more than the β_{\max} value of its parent node N' then there is no need to search below node N.

Alpha-Beta Analysis

- Pruning does not affect the final result.
- Assume a fixed branching factor and a fixed depth
- Best case: $b^{d/2} + b^{(d/2)-1}$
- Approximate as $b^{d/2}$
- Impact ?

Minmax: $10^9 = 1,000,000,000$

Alpha-beta: $10^5 + 10^4 = 110,000$

- But best-case analysis depends on choosing the best move first at cut nodes (not always possible)
- The worst case : No cut-offs, and Alpha-Beta degrades to Minmax



Summary

- Knowledge based agent (KBA)
 - Knowledge base and representation
- Inference Engine
 - Forward Chaining
 - Backward Chaining
- Game
 - MiniMax Algorithm
 - Alpha-beta Pruning