

Homework Assignment #2

(Due: Saturday, 25-Nov-2023, 11:59pm. Submit via Blackboard)

Question A. (35 marks)

- 1) What is a block anchor? How is it used in the implementation of a primary index? [4 marks]
- 2) Why can we have at most one primary or clustering index on a file, but several secondary indexes? [4 marks]
- 3) Consider the following *Movie* relation:

Movie (MovieID, Title, CategoryID, DirectorID, Rating)

MovieID is the primary key of the *Movie* relation. The *CategoryID* field contains 10 distinct values, and the records are evenly distributed among these values. The *DirectorID* field contains 200 distinct values, and the records are evenly distributed among these values. The *Movie* relation contains $N = 4,000$ records, with each record occupying 100 bytes. The file is sorted by the *CategoryID* field, and stored on a disk with the following configuration:

- Block size = 600 bytes
- Block pointer size = 6 bytes

Three indexes have been created on the *Movie* relation:

- A *clustering index* on the *CategoryID* field (Integer, 4 bytes)
- A *secondary index* on the *DirectorID* field (Integer, 4 bytes)
- A *bitmap index* on both *CategoryID* and *DirectorID* fields (1 byte = 8 bits; each bitmap is stored as a *fixed-length record*)

Consider the following SQL query:

```
SELECT * FROM Movie WHERE CategoryID = 3 AND DirectorID = 16;
```

Assume that 4 records are generated as the result of this SQL query. Which index is the most efficient for answering this query? Justify your answer by comparing the estimated query cost (i.e., the number of block accesses). [27 marks]

Answers for Question A. (35 marks)

- 1) Block anchor: First record in each data block [2 marks]
Primary index: One index entry for each data block, with the block anchor as the key [2 marks]
- 2) The primary or clustering index is built for sorted files. Each file can be sorted by only one column [2 marks]
The secondary index is built for unsorted files [2 marks]

3) Clustering index: [9 marks]

Index entry size = $4 + 6 = 10$ bytes

Index blocking factor = $\text{block size} / \text{index entry size} = 600 / 10 = 60$ entries/block

index entries = # distinct values = 10 entries

index blocks = $\text{\# index entries} / \text{index blocking factor} = \text{ceiling}(10 / 60) = 1$ block

Index access = 1 block

Each distinct *CategoryID* value has $4,000 / 10 = 400$ records

Record blocking factor = $\text{block size} / \text{record size} = 600 / 100 = 6$ records/block

Data access = $\text{ceiling}(400 / 6) = 67$ blocks

Total cost = index access + data access = $1 + 67 = 68$ blocks

Secondary index: [11 marks]

Index entry size = $4 + 6 = 10$ bytes

Index blocking factor = $\text{block size} / \text{index entry size} = 600 / 10 = 60$ entries/block

1st-level index entries = # distinct values = 200 entries

1st-level index blocks = $\text{\# 1st-level index entries} / \text{index blocking factor} = \text{ceiling}(200 / 60) = 4$ blocks

1st-level index access = $\log_2 B = \log_2 4 = 2$ blocks

Each block stores $600 / 6 = 100$ record pointers

Each distinct *DirectorID* value has $4,000 / 200 = 20$ records

Each distinct *DirectorID* value has $\text{ceiling}(20 / 100) = 1$ block of record pointers

2nd-level index access = 1 block

Data access = retrieve 20 records = 20 blocks (maximum)

Total cost = index access + data access = $2 + 1 + 20 = 23$ blocks

Bitmap index: [6 marks]

Bitmap size = # records = 4,000 bits = 500 bytes

blocks to store each bitmap = $\text{bitmap size} / \text{block size} = \text{ceiling}(500 / 600) = 1$ block

Index access to retrieve a *CategoryID* bitmap = 1 block

Index access to retrieve a *DirectorID* bitmap = 1 block

Data access = retrieve 4 final records = 4 blocks (maximum)

Total cost = index access + data access = $1 + 1 + 4 = 6$ blocks

Hence, bitmap index is the most efficient for answering this SQL query [1 mark]

Question B. (35 marks)

- 1) Consider two relations $R(A, B)$ and $S(\underline{A}, C)$. Field A is the primary key of relation S , and the foreign key of relation R referencing S . Assume that *page-oriented nested-loop join* is used for $R * S$ (natural join). Let R be the outer relation, and S be the inner relation. Given the following tuples of R and S , what are the first 4 tuples produced by $R * S$? Assume each block can store only 3 R tuples or S tuples. [8 marks]

- 1st tuple:
- 2nd tuple:
- 3rd tuple:
- 4th tuple:

A	B
7	x
2	z
9	y
8	y
3	w
9	x
1	w
3	y
5	z

R

A	C
8	1
4	3
2	6
1	5
3	7
5	1
7	8
9	2

S

- 2) Assume that R and S are stored on a disk with block size = 1000 bytes. Relation R contains 100,000 records, with each record occupying 20 bytes. Relation S contains 50,000 records, with each record occupying 50 bytes. Estimate the number of block accesses required for $R * S$ using page-oriented nested-loop join. [6 marks]
- 3) Can we improve the efficiency of $R * S$ by changing the join order, i.e. S as the outer relation and R as the inner relation? Justify your answer. [3 marks]
- 4) Consider the following relations in a *Company* database:

Employee (EID, EName, Email, Age, Address, Salary)
 Department (DID, DName)
 Join (EID, DID, Year)

EID is the primary key of *Employee*, and *DID* is the primary key of *Department*. (*EID*, *DID*) is the primary key of *Join*, where *EID* is the foreign key referencing *Employee* and *DID* is the foreign key referencing *Department*. *DName* is a unique field in *Department* relation. Given the following SQL query:

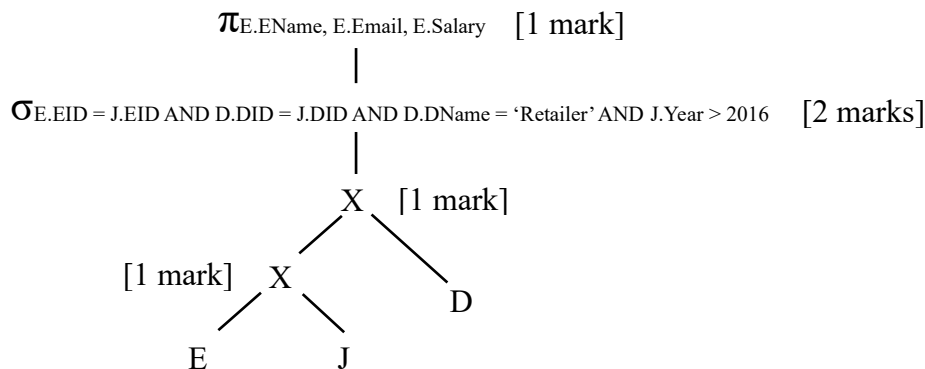
```
SELECT E.EName, E.Email, E.Salary
FROM Employee E, Join J, Department D
WHERE E.EID = J.EID AND D.DID = J.DID
      AND D.DName = 'Retailer' AND J.Year > 2016;
```

Show the initial query tree generated by the *conceptual evaluation strategy*. [5 marks]

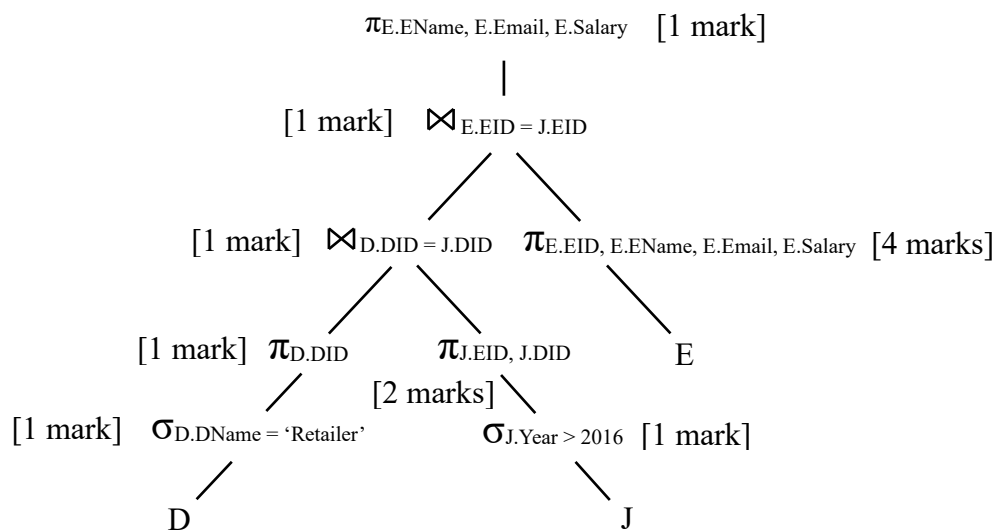
- 5) Show the most efficient query tree after applying all the five steps of heuristic-based *query tree optimization*. [13 marks]

Answers for Question B. (35 marks)

- 1) 1st tuple: (2, z, 6) [2 marks]
2nd tuple: (7, x, 8) [2 marks]
3rd tuple: (9, y, 2) [2 marks]
4th tuple: (8, y, 1) [2 marks]
- 2) Blocking factor of R = block size / record size = 1000 / 20 = 50 records/block [1 mark]
Blocking factor of S = block size / record size = 1000 / 50 = 20 records/block [1 mark]
 $B_R = \# \text{ records} / \text{block factor} = 100,000 / 50 = 2000 \text{ blocks}$ [1 mark]
 $B_S = \# \text{ records} / \text{block factor} = 50,000 / 20 = 2500 \text{ blocks}$ [1 mark]
 $\# \text{ block accesses} = B_R + B_R * B_S = 2000 + 2000 * 2500 = 5,002,000 \text{ blocks}$ [2 marks]
- 3) $\# \text{ block accesses} = B_S + B_S * B_R = 2500 + 2500 * 2000 = 5,002,500 \text{ blocks}$ [2 marks]
No, we cannot improve the efficiency of R * S by changing the join order, i.e., I/O cost is increased [1 mark]
- 4) Initial query tree:



- 5) Optimized query tree:**



[1 mark] Change the order of E and D

Question C. (30 marks)

- 1) *Two-phase locking* (2PL) is widely used for concurrency control in a relational database system. Consider the following schedule for two transactions T1 and T2:

S: W1 (X) ; R2 (Y) ; R1 (Y) ; R2 (X) ; C1 ; C2

“C1” means transaction T1 commits. For each of the following 2PL protocols: (1) Briefly describe the protocol; (2) State if the protocol allows schedule S, i.e., allows the actions to occur in exactly the order shown in schedule S; (3) Explain the reason why schedule S is allowed or not allowed under the protocol. [15 marks]

- Basic 2PL
- Conservative 2PL
- Strict 2PL

- 2) Consider two transactions T1, T2 and two data items X, Y.

T1: R (X) ; W (X) ; R (Y) ; W (Y)

T2: R (X) ; W (X)

Given the following schedule of interleaved operations from T1 and T2:

S: R1 (X) ; W1 (X) ; R2 (X) ; R1 (Y) ; W2 (X) ; C2 ; W1 (Y) ; A1

“A1” means transaction T1 aborts, and “C2” means transaction T2 commits. Answer each of the following questions. [15 marks]

- What type of *anomaly* occurs in this schedule S? Explain your answer.
- For each of T1 and T2, insert all the lock and unlock operations to make the transaction satisfy the strict 2PL protocol.
- Show that modifying the schedule according to strict 2PL can prevent the anomaly. Justify your answer.

Answers for Question C. (30 marks)

1) Basic 2PL:

Growing phase (Whenever the scheduler receives an operation on any item, it must acquire a lock on that item before execution. No locks can be released in this phase) [1 mark] +

Shrinking phase (Once the scheduler has released a lock for a transaction, it cannot request any additional locks on any data item for this transaction) [1 mark]

S is allowed [1 mark]

Reason: The write lock on X can be released by T1 after R1(Y), T1 is in the shrinking phase; Then T2 can obtain a read lock on X for R2(X), T2 is in the growing phase [2 marks]

Conservative 2PL:

Growing phase + shrinking phase [1 mark] Lock all desired data items before transaction starts [1 mark]

S is allowed [1 mark]

Reason: The write lock on X can be released by T1 after W1(X), T1 is in the shrinking phase; Then T2 can obtain read locks on both X and Y to start from R2(Y), T2 is in the growing phase [2 marks]

Strict 2PL:

Growing phase + shrinking phase [1 mark] Write locks cannot be released until transaction terminates [1 mark]

S is not allowed [1 mark]

Reason: The write lock on X will be released by T1 only at commit time C1. Hence, T2 cannot obtain the read lock on X to run R2(X) [2 marks]

2) Anomaly: Dirty read [2 marks]

Reason: T2 uses item X which has been modified by the aborted transaction T1 [2 marks]

T1: WL(X); R(X); W(X); WL(Y); R(Y); W(Y); UL(X); UL(Y) [4 marks]

T2: WL(X); R(X); W(X); UL(X) [2 marks]

S': R1(X); W1(X); R1(Y); W1(Y); A1; R2(X); W2(X); C2 [3 marks]

Reason: T2 uses data item X only after T1 has already been aborted [2 marks]